

# Deep Learning: Algorithms and Applications

Peking University

Spring 2023

Prof. Yadong Mu

---

## Homework 1

Due Date: May 15th

### 1. MLP for Classification

In a multi-class classification problem, *e.g.* a problem with  **$M$  classes**, we use a model  $f_\theta$  to produce the output  $\hat{y}$ .  $\hat{y}$  is a vector of dimension  $M$  with sum 1, whose components indicate the probability of  $x$  belong to each class.

In this section, you will use a simple model — the 3-layer MLP — to tackle this problem. Assume that we have  $K$  labeled data  $(x_k, y_k)$  with  $k = 1, 2, \dots, K$  where  $x_k \in \mathbb{R}^N$  and  $y_k \in \mathbb{R}^M$  is a  $M$ -dimensional one-hot vector.

- Define the cross entropy loss for a multi-class classification problem.**
- A 3-layer MLP consists of an input layer, a hidden layer and an output layer with two learned parameter matrices  $W^1$ ,  $W^2$  between successive layers. Please note that there is generally a Softmax layer after the output layer to scale the output, which we omitted in this sub-question for simplicity. When given an input  $x$ , this model performs the following forward computations sequentially:

$$\begin{aligned}a_1 &= W^1 x, \\h &= \sigma(a_1), \\a_2 &= W^2 h, \\\hat{y} &= \sigma(a_2).\end{aligned}$$

where  $W^1 \in \mathbb{R}^{D \times N}$ ,  $W^2 \in \mathbb{R}^{M \times D}$  are parameter matrices and  $\sigma(z) = \frac{1}{1+e^{-z}}$  is the element-wise Sigmoid function. **Use the loss function you defined in sub-question a), apply an SGD update to the parameters  $W^1$  and  $W^2$  with learning rate  $\eta$  via back-propagation.** You must identify the **closed-form gradient** of each parameter, which is  $\frac{\partial \mathcal{L}}{\partial W^2(m,d)}$  and  $\frac{\partial \mathcal{L}}{\partial W^1(d,n)}$  in your derivation.

[HINT: 1) For simplicity, you can assume that the SGD program uses only one training data per batch. 2) Use formula  $\sigma'(z) = \sigma(z)(1 - \sigma(z))$  for a scalar  $z$ . ]

### 2. Dropout and Regularization

Dropout is a well-known way to prevent neural networks from overfitting. In this section, you will show this regularization explicitly for linear regression. Recall that linear regression optimizes  $w \in \mathbb{R}^d$  to minimize the following MSE objective:

$$\mathcal{L}(w) = \|y - Xw\|^2$$

where  $y \in \mathbb{R}^n$  is the response to design matrix  $X \in \mathbb{R}^{n \times d}$ . One way of using dropout during training on the  $d$ -dimensional input features  $x_i$  involves *keeping* each feature at random with probability  $p$  (and zero out it if not kept).

- a) **Show that when we apply such dropout, the learning objective becomes**

$$\mathcal{L}(w) = \mathbb{E}_{M \sim \text{Bernoulli}(p)} \|y - (M \odot X)w\|^2$$

where  $\odot$  denotes the element-wise product and  $M \in \{0, 1\}^{n \times d}$  is a random mask matrix whose element  $m_{i,j}$  have *i.i.d.* Bernoulli distribution with success probability  $p$ .

- b) **Show that we can manipulate the dropout learning objective to a explicit regularized objective:**

$$\mathcal{L}(w) = \|y - pXw\|^2 + p(1 - p)\|\Gamma w\|^2$$

**and define a suitable matrix  $\Gamma$ .**